

Name: Thomas Hoang

Student ID: A1984077

Assignment: Assignment 2.2 Programming Task 2

Project: NumCompute Stream

Overview

This project is a streaming version of NumCompute. Instead of training once on all rows, the data is fed in small chunks and the pipeline updates step by step.

Requirement	Project part
Streaming learning	Models, preprocessing, stats, and metrics support partial_fit() or update()
Tree model	DecisionTreeClassifier
Ensemble model	BaggingClassifier with several trees
Visualisation	visualise.py with plots for accuracy and model comparison
Demo	demo/stream_demo.ipynb using the white wine CSV
Benchmark	Tree vs bagging under the same streaming chunks

Package Design

The new work is kept in **numcompute_stream/**. This keeps it separate from the old Assignment 2.1 code, so the batch version does not get mixed with the streaming one.

File	Purpose
stats.py	Keep running mean, variance, and histogram values
metrics.py	Keep accuracy, confusion matrix, precision, recall, and F1 over chunks
preprocessing.py	Update scaler, imputer, and encoder from new chunks
pipeline.py	Connect preprocessing and model with one shared partial_fit() call
tree.py	Grow one decision tree as more chunks arrive
ensemble.py	Train several trees and vote between them
stream.py	Train one chunk, score it, and save logs
visualise.py	Make plots for the demo and benchmark

API choice	Why it is used
partial_fit()	Used when a component learns from a new chunk
update()	Used when a metric or statistic adds new values
predict()	Kept the same for tree and bagging
StreamTrainer	Keeps the notebook shorter and stores training logs

Model Design

Model	How it works	Notes
DecisionTreeClassifier	One tree stores samples at leaves and tries to split when enough data is received	Fast and easy to inspect
BaggingClassifier	Several trees train on bootstrap samples from each chunk	Slower but usually more stable

Bagging was chosen because it is simple to connect with streaming chunks. For each new chunk, every tree gets its own sampled version of the chunk. After that, prediction is just majority voting.

Demo Setup

The demo notebook loads the white wine dataset with **numcompute.io.load_csv()**. The quality column is changed into a binary target. Quality 6 or higher becomes class 1, and the rest becomes class 0.

Demo item	Result
Raw data shape	4898 rows and 12 columns
Feature shape	4898 rows and 11 features
Class count	1640 class 0 and 3258 class 1
Chunk size	50 rows
Number of chunks	98
Models compared	Tree and bagging

Step	What happens
1	Load CSV
2	Split into chunks
3	Update StandardScaler
4	Train model with partial_fit()
5	Predict on the current chunk
6	Save accuracy in StreamTrainer.log__
7	Plot tree vs bagging accuracy

Demo Results

The notebook shows the model working chunk by chunk. The last chunk and final cumulative scores are below.

Model	Last chunk accuracy	Final cumulative accuracy
Tree	0.667	0.713
Bagging	0.729	0.705

The bagging model was also checked with streaming metrics after training.

Metric	Result
Precision	0.957
Recall	0.696
F1	0.806

Confusion matrix	Predicted 0	Predicted 1
Actual 0	280	1360
Actual 1	139	3119

The scaler also updated through the stream. The first three final mean values in the demo were **6.8548**, **0.2782**, and **0.3342**.

Stability and Edge Cases

This section shows the checks used so the code does not break when chunks are small, missing values appear, or a split is not useful.

Problem	Handling
Running mean and variance	Keep count, mean, and spread values, then merge each new chunk into them
Zero standard deviation	Replace zero scale with 1.0 before transform
NaN values	The imputer skips NaNs when updating the column mean
Empty split in tree	Reject the split
One-class leaf	Do not split it
Same feature values	Skip that feature for splitting
Tie between split options	Choose the lower feature and lower threshold
Wrong input shape	Raise a clear error before training

Testing

There are 81 tests for `numcompute_stream/`, which is more than the required 30.

Test area	What is checked
Stats	Chunk updates match expected combined results
Metrics	Accuracy and other metrics build over multiple updates
Preprocessing	Scaler, imputer, and encoder learn over chunks
Pipeline	Steps are called in the right order
Tree	Fit, predict, depth, single-class data, and split cases
Ensemble	Tree count, bootstrap updates, and voting
StreamTrainer	Logs, scoring, and cumulative accuracy
Visualise	Plot saving and invalid inputs

Testing helped me find small issues faster, especially shape mismatches and cases where an empty chunk should fail.

Benchmark Results

For benchmark, the base tree and the ensemble are compared on the same streaming chunks.

Benchmark item	Value
Dataset	assets/winequality-white.csv
Chunk size	50
Chunks	98
Base model	DecisionTreeClassifier
Ensemble model	BaggingClassifier with 5 trees

Model	Avg fit time per chunk	Avg predict time per chunk	Final chunk accuracy	Final cumulative accuracy
Tree	0.72 ms	0.03 ms	0.667	0.713
Bagging	2.73 ms	0.34 ms	0.729	0.705

From the result, bagging is slower because it trains and predicts with five trees. With this seed, bagging has better last chunk accuracy, but the tree has better cumulative accuracy.

Problems Faced

Problem	Fix
Notebook import errors	Added a setup cell to find the package root
Streaming tree design	Stored samples and class counts at leaves before splitting
Zero variance values	Made scaler avoid dividing by zero
Bagging randomness	Used random_state=42 for repeatable demo results
Too much notebook code	Moved training and logging into StreamTrainer

The tree part took the most time because a normal decision tree expects all data at once. The main question was how a leaf should grow when only a chunk arrives.

Reflection

This assignment helped me learn more about how decision trees are actually built. It was harder than expected, especially when the tree had to grow chunk by chunk instead of from one full dataset. AI was helpful for understanding the steps, but the code still needed checking and testing to make sure it matched the assignment.

Luckily, the old NumCompute codebase gave a good base to work from. The existing validation style, preprocessing ideas, and CSV loading made the streaming package easier to organise.

The hardest part was deciding how to grow the tree. A normal tree can see all data at once, but this version only receives one chunk at a time. The leaf needs to store enough information, update class counts, and only split when the data is useful.

There was also a lot of tedious work writing tests. The tests were not difficult one by one, but many edge cases had to be checked, like empty chunks, wrong shapes, zero variance, single-class data, and voting behaviour. This part took time, but it made the final package more reliable.

AI Use Acknowledgement

Generative AI was used for planning, guiding the structure, and explaining hard parts of the assignment.

It also helped with repetitive work, especially drafting test cases and checking report wording.

The code, tests, demo results, and benchmark output were still reviewed and verified by me. I am fully aware of how the submitted code works.